



## **ALY 6015: INTERMEDIATE ANALYTICS**

Assignment 4: Feature Selection in R

Submitted to

Prof. Fatemeh Ahmadi Abkenari

Submitted by

Abhilash Dikshit

Mrityunjay Gupta

Siddharth Alashi

Smit Parmar

## Assignment 4: Feature Selection in R

Abhilash Dikshit, Siddharth Alashi ,Mrityunjay Gupta, Smit Parmar  
*College of Professional Studies*  
*Northeastern University*  
*Vancouver, Canada*

### I. Abstract:

A built-in dataset in R called “mtcars” provides measurements for 32 distinct cars over 11 different attributes. In this paper we will be summarizing methods to optimize model using feature selection techniques. Forwards selection techniques and Both direction regression method helps us to select the best regression model on this dataset. The ‘mtcars’ dataset is split into Train and Test dataset with the ratio of 70/30. Furthermore, we have made visualizations and descriptive analysis which describes the comparison between the variables. The importance of the models functioning best amongst them is then summarized by ANOVA TEST. Finally, the references provides a support for our arguments, ideas, and opinions.

### II. Introduction

The data was extracted from the 1974 Motor Trend US magazine and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973–74 models)

This dataset shall consist of 11 columns and 32 observations which are labelled below.

mpg	Miles/(US) gallon
cyl	Number of cylinders
disp	Displacement (cu.in.)
hp	Gross horsepower
drat	Rear axle ratio
wt	Weight (1000 lbs)
qsec	1/4 mile time
vs	Engine (0 = V-shaped, 1 = straight)
am	Transmission (0 = automatic, 1 = manual)

gear	Number of forward gears
carb	Number of carburetors

It is frequently required and beneficial to divide the data set into training and testing sets. The model will be trained on the training set of data, and the test set will be used to evaluate the model. This makes sure that we are not overfitting the model and that it functions properly with new data. It is extremely usual to utilise a 70/30 split, where 70% of the observations are used for the training set and 30% are used for testing. Fig1. and Fig2. describes the structure of Train and Test dataset

```
> str(sample_train)
'data.frame': 22 obs. of 11 variables:
 $ mpg : num 21 21.4 18.7 18.1 22.8 17.8 16.4 17.3 15.2 14.
 $ cyl : num 6 6 8 6 4 6 8 8 8 8 ...
 $ disp: num 160 258 360 225 141 ...
 $ hp : num 110 110 175 105 95 123 180 180 180 230 ...
 $ drat: num 3.9 3.08 3.15 2.76 3.92 3.92 3.07 3.07 3.07 3.
 $ wt : num 2.88 3.21 3.44 3.46 3.15 ...
 $ qsec: num 17 19.4 17 20.2 22.9 ...
 $ vs : num 0 1 0 1 1 1 0 0 0 0 ...
 $ am : num 1 0 0 0 0 0 0 0 0 0 ...
 $ gear: num 4 3 3 3 4 4 3 3 3 3 ...
 $ carb: num 4 1 2 1 2 4 3 3 3 4 ...
```

Fig1. Structure of Train dataset

```
> str(sample_test)
'data.frame': 10 obs. of 11 variables:
 $ mpg : num 21 22.8 14.3 24.4 19.2 10.4 10.4 32.4 33.9 1
 $ cyl : num 6 4 8 4 6 8 8 4 4 8
 $ disp: num 160 108 360 147 168 ...
 $ hp : num 110 93 245 62 123 205 215 66 65 335
 $ drat: num 3.9 3.85 3.21 3.69 3.92 2.93 3 4.08 4.22 3.5
 $ wt : num 2.62 2.32 3.57 3.19 3.44 ...
 $ qsec: num 16.5 18.6 15.8 20 18.3 ...
 $ vs : num 0 1 0 1 1 0 0 1 1 0
 $ am : num 1 1 0 0 0 0 0 1 1 1
 $ gear: num 4 4 3 4 4 3 3 4 4 5
 $ carb: num 4 1 4 2 4 4 4 1 1 8
```

Fig2. Structure of Test dataset

### III. Descriptive Analysis

The summary in Fig3. Concludes the introduction of the statistics to 9 variables in the dataset.

1. we find that the average miles per gallon for 22 cars is 19.96mpg. 25% of cars have 15.96 mpg as their average and 75% of the cars have 21.48mpg as their average.
2. The average number of cylinders among 22 cars is 6 cylinder but the maximum number of cylinders in the cars are 8 cylinder cars.
3. The average horse-power (hp) of the cars is 144.3 hp, whereas the minimum horsepower of the car is only 52hp and the maximum horsepower is 264hp.
4. Average weight (wt) of cars is 3.16 tons. the heaviest cars are of 5.34 tons.

```
> summary(sample_train)
```

mpg	cyl	disp
Min. :13.30	Min. :4.000	Min. : 75.7
1st Qu.:15.95	1st Qu.:4.000	1st Qu.:126.0
Median :18.95	Median :6.000	Median :241.5
Mean :19.96	Mean :6.273	Mean :229.9
3rd Qu.:21.48	3rd Qu.:8.000	3rd Qu.:314.5
Max. :30.40	Max. :8.000	Max. :440.0

hp	drat	wt
Min. : 52.0	Min. :2.76	Min. :1.513
1st Qu.:106.0	1st Qu.:3.08	1st Qu.:2.772
Median :136.5	Median :3.66	Median :3.325
Mean :144.3	Mean :3.58	Mean :3.161
3rd Qu.:178.8	3rd Qu.:3.92	3rd Qu.:3.678
Max. :264.0	Max. :4.93	Max. :5.345

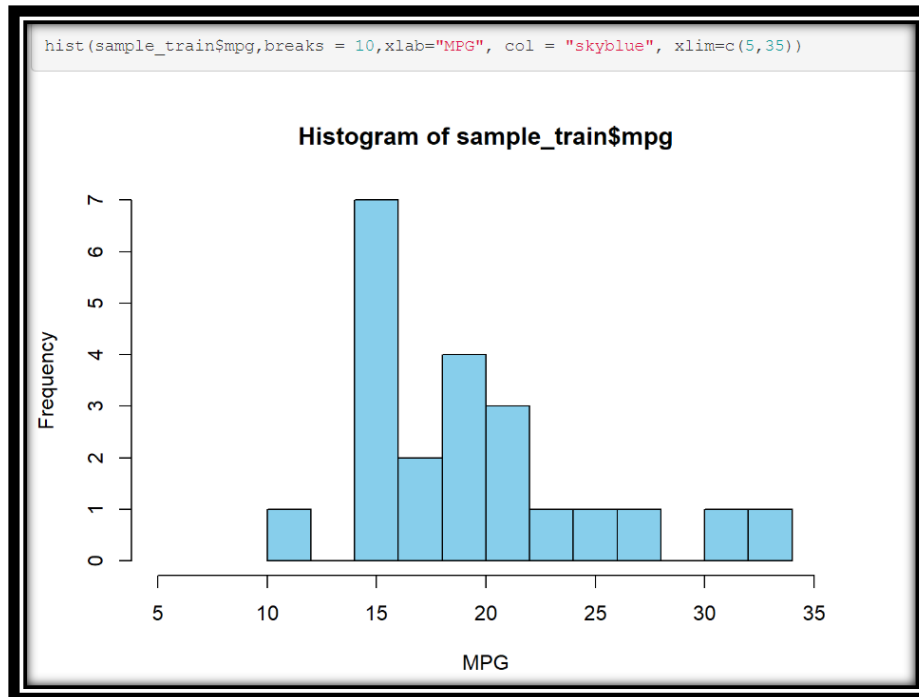
  

qsec	vs	am
Min. :14.50	Min. :0.0000	Min. :0.0000
1st Qu.:16.93	1st Qu.:0.0000	1st Qu.:0.0000
Median :17.41	Median :0.0000	Median :0.0000
Mean :17.83	Mean :0.4091	Mean :0.3636
3rd Qu.:18.82	3rd Qu.:1.0000	3rd Qu.:1.0000

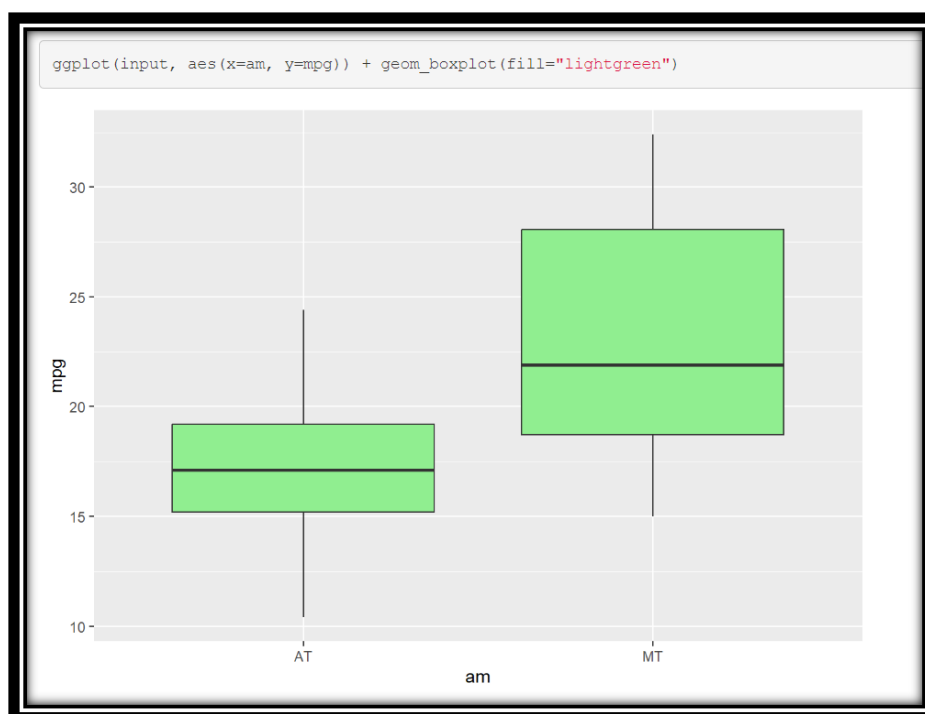
*Fig3. Summary of 'mtcars' dataset*

### III. Exploratory Analysis

1. The distribution of the outcome variable (mpg) is plotted using a histogram which suggests a resemblance with normal distribution. Furthermore, the maximum number of cars has 15 mpg as their average, almost 45% of the cars have an average more than the Mean calculated.

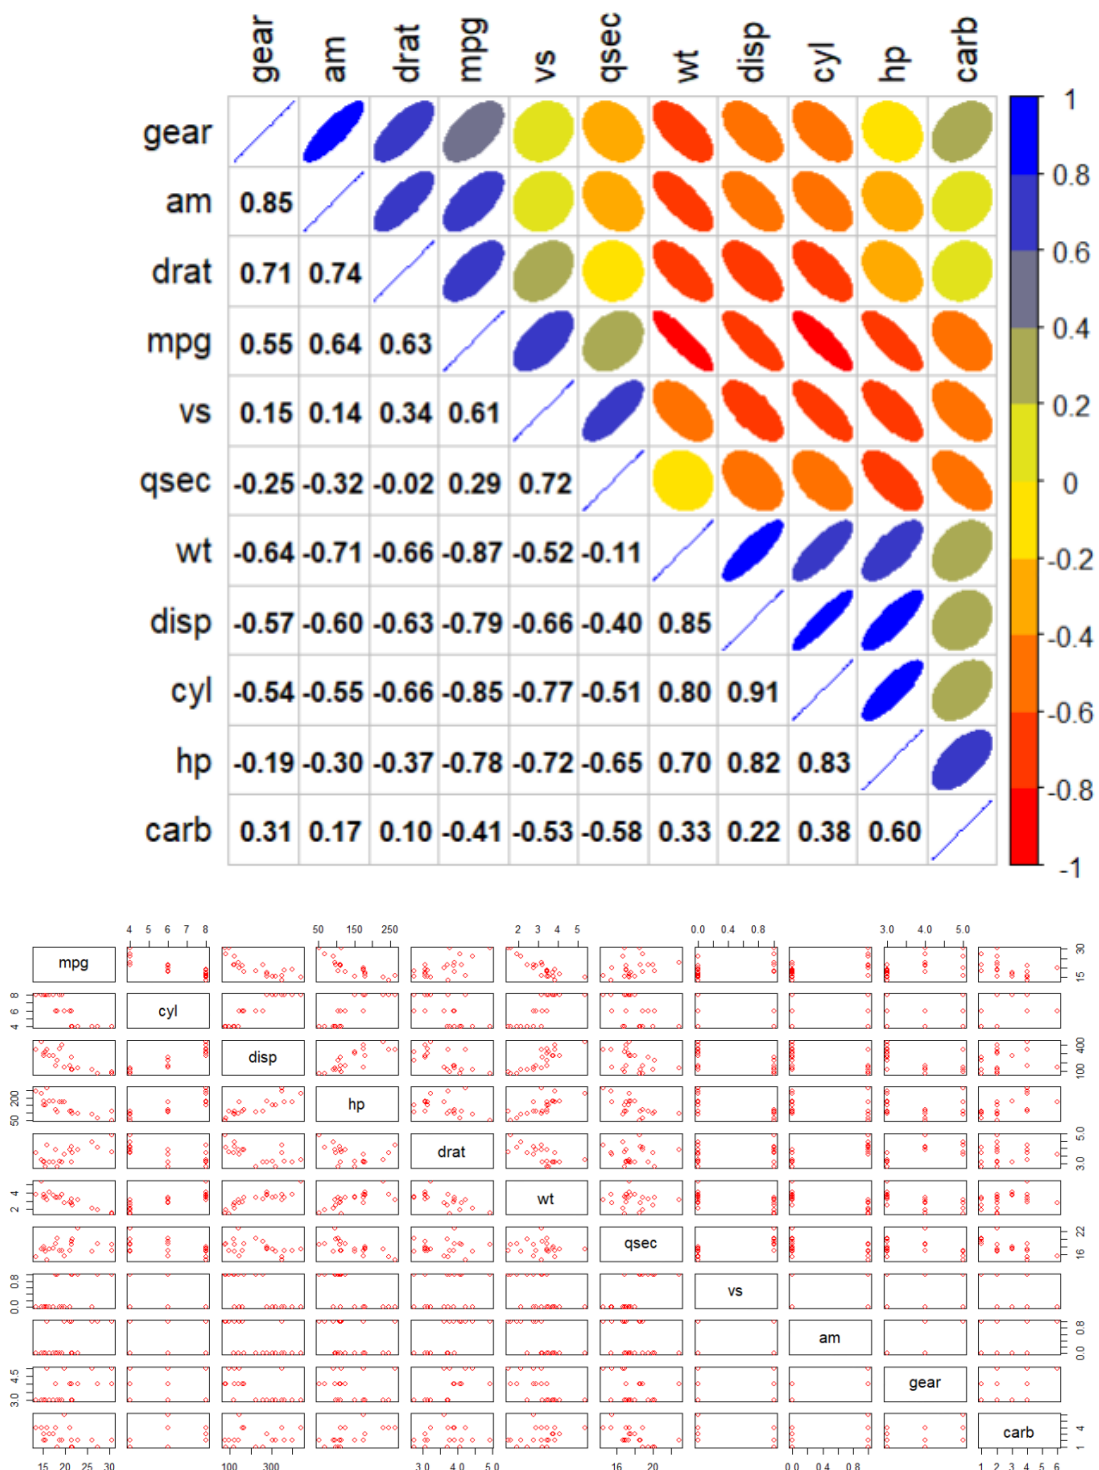


2. A boxplot of the outcome variable (mpg) is plotted with (am). It suggests manual transmission is better for mpg as compared to an automatic transmission.



3. To check the co-linearity between the variables a **Pair plot** is plotted. The Pair plot shows a strong relationship between different variables and miles per gallon. we can conclude from the Fig6. and Fig7.

- Gear has strong positive linear relationship between Transmission , real axel ratio and Negative weakly linear relationship with weight, displacement, Cylinder, and horsepower.
- Transmission has weak positive linear relationship with carburetors. Also, it has weak negative linear relationship with qsec, horsepower(hp), cylinder(cyl).
- Miles per gallon (mpg) has strong positive linear co-relationship with Engine(vs). Whereas, it has Strong Negative Linear relationship with weight(wt), displacement(dis), cylinder(cyl), horsepower(hp).



## Feature Selection Method

By adding and removing predictors from the model progressively until there is no longer a statistically legitimate reason to add or remove any more, stepwise regression is a technique we may use to create a regression model from a set of predictor variables.

With this model we have used Stepwise regression as

- Forward Selection
- Both-Direction Selection.

### Forward selection

The first method is the forward selection method. In this case, we start with no predictors and then add the predictor with the highest correlation with the response variable.

By including the variable, we ensure that the model has actually improved.

If it has, repeat the process. When there are no more improvements that can be made by adding variables to the model, the process will end. By setting the 'direction' parameter to "forward," we select the step() function's forward selection method.

```
step(lm(mpg ~ 1, data = mtcars), direction = 'forward', scope = ~ disp + hp + drat + wt + qsec)
```

```
## Start:  AIC=115.94
## mpg ~ 1
##
##      Df Sum of Sq  RSS   AIC
## + wt    1   847.73 278.32  73.217
## + disp   1   808.89 317.16  77.397
## + hp     1   678.37 447.67  88.427
## + drat   1   522.48 603.57  97.988
## + qsec   1   197.39 928.66 111.776
## <none>          1126.05 115.943
##
## Step:  AIC=73.22
## mpg ~ wt
##
##      Df Sum of Sq  RSS   AIC
## + hp     1    83.274 195.05  63.840
## + qsec   1    82.858 195.46  63.908
## + disp   1    31.639 246.68  71.356
## <none>          278.32  73.217
## + drat   1     9.081 269.24  74.156
##
## Step:  AIC=63.84
## mpg ~ wt + hp
##
##      Df Sum of Sq  RSS   AIC
## <none>          195.05  63.840
## + drat   1    11.3659 183.68  63.919
## + qsec   1     8.9885 186.06  64.331
## + disp   1     0.0571 194.99  65.831
```

```
##
## Call:
## lm(formula = mpg ~ wt + hp, data = mtcars)
##
## Coefficients:
## (Intercept)          wt          hp
##    37.22727    -3.87783    -0.03177
```

Here we can see that after applying the Forward Selection method we found the best model

$$mpg \sim wt + hp$$

Here the above conclusion comes up based on their AIC values which are minimum(minimum AIC gives the best model) for the above equation.

When we performed Linear regression of the above equation with a given data set our output is justified

#### Linear regression of Forward Selection Methods Result

```
model_forward <- lm(formula = mpg ~ wt + hp, data = mtcars)
summary(model_forward)
```

```
##
## Call:
## lm(formula = mpg ~ wt + hp, data = mtcars)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.941 -1.600 -0.182  1.050  5.854
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  37.22727    1.59879   23.285 < 0.0000000000000002 ***
## wt          -3.87783    0.63273   -6.129  0.00000112 ***
## hp          -0.03177    0.00903   -3.519  0.00145 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.593 on 29 degrees of freedom
## Multiple R-squared:  0.8268, Adjusted R-squared:  0.8148
## F-statistic: 69.21 on 2 and 29 DF, p-value: 0.0000000000009109
```

Here we have performed Linear regression validating our output of the forward selection method as we can see the important predictors are identified in the output with intercept



The following code and the output pasted show the performance of the Both-Direction stepwise selection method.

```
options(scipen = 100)
model_step <- step(lm(mpg ~ ., data = mtcars), direction = 'both')
summary(model_step)
```

```
Start: AIC=70.9
mpg ~ cyl + disp + hp + drat + wt + qsec + vs + am + gear
```

	Df	Sum of Sq	RSS	AIC
~ cyl	1	0.0799	147.57	68.915
~ vs	1	0.1601	147.66	68.932
~ carb	1	0.4067	147.90	68.986
~ gear	1	1.3531	148.85	69.190
~ drat	1	1.6270	149.12	69.249
~ disp	1	3.9167	151.41	69.736
~ hp	1	6.8399	154.33	70.348
~ qsec	1	8.8641	156.36	70.765
~ none			147.49	70.898
~ am	1	10.5467	158.04	71.108
~ wt	1	27.0144	174.51	74.280

```
Step: AIC=68.92
mpg ~ disp + hp + drat + wt + qsec + vs + am + gear + carb
```

	Df	Sum of Sq	RSS	AIC
- vs	1	0.2685	147.84	66.973
- carb	1	0.5201	148.09	67.028
- gear	1	1.8211	149.40	67.308
- drat	1	1.9826	149.56	67.342
- disp	1	3.9009	151.47	67.750
- hp	1	7.3632	154.94	68.473
<none>			147.57	68.915
+ qsec	1	10.0933	157.67	69.032
+ am	1	11.8359	159.41	69.384
+ wt	1	0.0799	147.49	70.898
+ cyl	1	27.0280	174.60	72.297

```
Step: AIC=66.97
mpg ~ disp + hp + drat + wt + qsec + am + gear + carb
```

	Df	Sum of Sq	RSS	AIC
- carb	1	0.6855	148.53	65.121
- gear	1	2.1437	149.99	65.434
- drat	1	2.2139	150.06	65.449
- disp	1	3.6467	151.49	65.753
- hp	1	7.1060	151.95	66.475
<none>			147.84	66.973
+ am	1	11.5694	159.41	67.384
+ qsec	1	15.6830	163.53	68.200
+ vs	1	0.2685	147.57	68.915
+ cyl	1	0.1883	147.66	68.932
+ wt	1	27.3799	175.22	70.410

```
Call:
lm(formula = mpg ~ wt + qsec + am, data = mtcars)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-3.4811 -1.5555 -0.7257  1.4110  4.6610
```

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	9.6178	6.9596	1.382	0.177915	
wt	-3.9165	0.7112	-5.507	0.0000695	***
qsec	1.2259	0.2887	4.247	0.000216	***
am	2.9358	1.4109	2.081	0.046716	*

```
Signif. codes:
0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 2.459 on 28 degrees of freedom  
Multiple R-squared: 0.8497, Adjusted R-squared: 0.8336  
F-statistic: 52.75 on 3 and 28 DF, p-value: 0.0000000000121

```
Step: AIC=65.12
mpg ~ disp + hp + drat + wt + qsec + am + gear
```

	Df	Sum of Sq	RSS	AIC
- gear	1	1.565	150.09	63.457
- drat	1	1.932	150.46	63.535
<none>			148.53	65.121
- disp	1	10.110	158.64	65.229
- am	1	12.323	160.85	65.672
- hp	1	14.826	163.35	66.166
+ carb	1	0.685	147.84	66.973
+ vs	1	0.434	148.09	67.028
+ cyl	1	0.414	148.11	67.032
- qsec	1	26.408	174.94	68.358
- wt	1	69.127	217.66	75.350

```
Step: AIC=63.46
mpg ~ disp + hp + drat + wt + qsec + am
```

	Df	Sum of Sq	RSS	AIC
- drat	1	3.345	153.44	62.162
- disp	1	8.545	158.64	63.229
<none>			150.09	63.457
- hp	1	13.285	163.38	64.171
+ gear	1	1.565	148.53	65.121
+ cyl	1	1.003	149.09	65.242
+ vs	1	0.645	149.45	65.319
+ carb	1	0.107	149.99	65.434
- am	1	20.036	170.13	65.466
- qsec	1	25.574	175.67	66.491
- wt	1	67.572	217.66	73.351

```
Step: AIC=62.16
mpg ~ disp + hp + wt + qsec + am
```

	Df	Sum of Sq	RSS	AIC
- disp	1	6.629	160.07	61.515
<none>			153.44	62.162
+ hp	1	12.572	166.01	62.682
+ drat	1	3.345	150.09	63.457
+ gear	1	2.977	150.46	63.535
+ cyl	1	2.447	150.99	63.648
+ vs	1	1.121	152.32	63.927
+ carb	1	0.011	153.43	64.160
- qsec	1	26.470	179.91	65.255
- am	1	32.198	185.63	66.258
- wt	1	68.843	233.48	72.051

The procedure information for the **Both-direction** Stepwise regression is:

As with the forward-stepwise selection, we added predictors to the model successively. After including each predictor, we did, however, also delete any predictors that were no longer improving the model's fit.

The final model turns out to be:

$$\text{mpg} \sim 9.62 - 3.92 \cdot \text{wt} + 1.23 \cdot \text{qsec} + 2.94 \cdot \text{am}$$

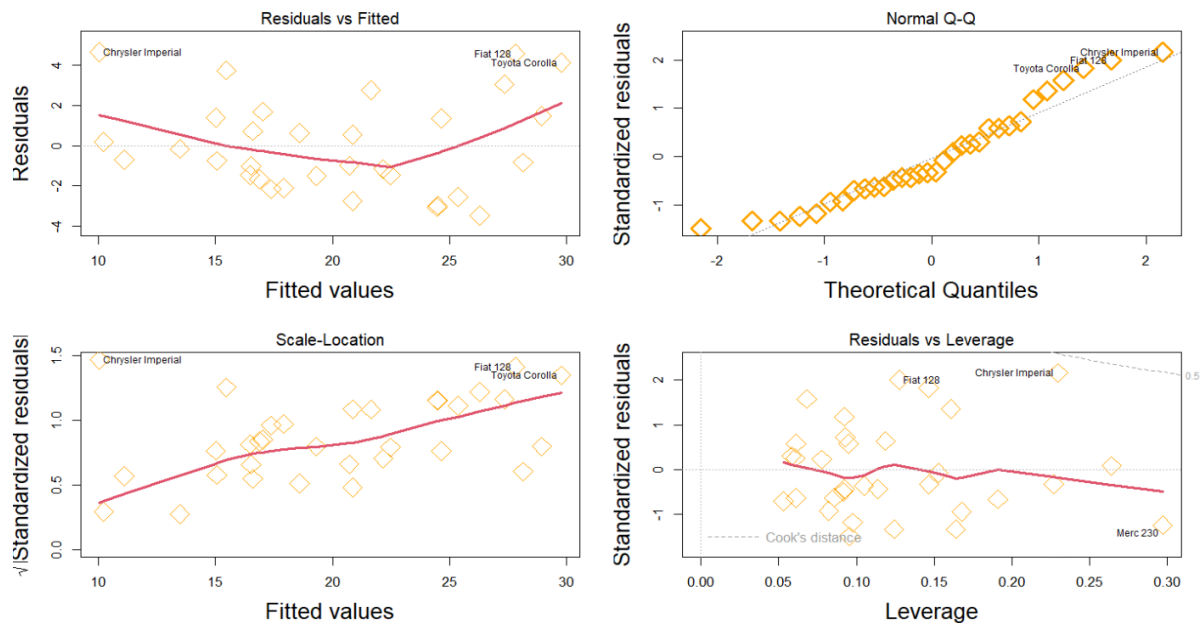
**The formula describes that with 1% change such as increase in miles per gallon (mpg) will result in -3.92 % decrease in weight and 1.23% increase in qsec , lastly 2.94% increase in Transmission.**

### **Residual Plots and Diagnostics**

Plot analysis from left to right in :

- 1) The residuals, distance of a point to the regression line, do not show a pattern as they have a random scatter about the dotted line.
- 2) The residuals in the Quantile/Quantile plot for the most part follow the line and can be assumed to be normally distributed,
- 3) The red line is fairly flat demonstrating homoschedasity, the residuals are not affected by explanatory variables
- 4) None of the residuals have a Cook's distance of greater than 0.5.

**In conclusion,** the type of car transmission that achieves better fuel efficiency is uncertain as other car attributes; horsepower, car weight and number of cylinders, may be a better indication of fuel efficiency. This model could be further refined through such techniques such as reducing any covariance between variables such as horsepower and number of cylinders or weight.



## Model Comparison

We are performing the model comparison of the Results of the Forward Selection method and Stepwise Selection Methods to determine which method provides the better Selection.

We are performing three comparison methods namely ANOVA, AIC, and BIC

### Compare Models With Anova

```
fit1 <- lm(formula = mpg ~ wt, data = mtcars)
fit2 <- lm(formula = mpg ~ wt + hp, data = mtcars)
anova(fit1, fit2)
```

	Res.Df <dbl>	RSS <dbl>	Df <dbl>	Sum of Sq <dbl>	F <dbl>	Pr(>F) <dbl>
1	30	278.3219	NA	NA	NA	NA
2	29	195.0478	1	83.27418	12.38133	0.001451229

2 rows

The Anova value of Fit 1(forward Selection) has the Anova null whereas the value of Anova for Fit 2 is 0.014 so, clearly, we can see fit 2 gives the best result than model 1 stating that the forward selection method is accurate with the given dataset.

### Compare models with AIC

```
AIC(fit1, fit2)
```

	df <dbl>	AIC <dbl>
fit1	3	166.0294
fit2	4	156.6523

2 rows

The AIC value of Fit 1 is 166.0294 and the AIC value of Fit is 156.6523 So, we can say that compared models using AIC methods which establish our result about fit 2 or forward selection was the best for the given dataset

#### Compare models with BIC

### Compare models with BIC

```
BIC(fit1, fit2)
```

	df <dbl>	BIC <dbl>
fit1	3	170.4266
fit2	4	162.5153

2 rows

We have performed model testing By using BIC and the obtained value of fit 1 is 170.4266 and the fit 2 value is 162.5193 So, clearly, we got the output validating our previous two testing methods stating that fit 2 or forward selection was perfect to give an accurate model for provided dataset

## Dataset -2 Hitters Dataset.

This Hitters data collection was obtained via the Carnegie Mellon University-maintained StatLib library. This is a portion of the data that was used in the poster session for the 1988 ASA Graphics Section. The pay information was first published in Sports Illustrated on April 20, 1987. The 1987 Baseball Encyclopedia Update, published by Collier Books, Macmillan Publishing Company, New York, provided the 1986 and career statistics.

<b>AtBat</b>	Number of times at bat in 1986
<b>Hits</b>	Number of hits in 1986
<b>HmRun</b>	Number of home runs in 1986
<b>Runs</b>	Number of runs in 1986
<b>RBI</b>	Number of runs batted in in 1986
<b>Walks</b>	Number of walks in 1986
<b>Years</b>	Number of years in the major leagues
<b>CAtBat</b>	Number of times at bat during his career
<b>CHits</b>	Number of hits during his career
<b>CHmRun</b>	Number of home runs during his career
<b>CRuns</b>	Number of runs during his career
<b>CRBI</b>	Number of runs batted in during his career
<b>CWalks</b>	Number of walks during his career
<b>League</b>	A factor with levels A and N indicating player's league at the end of 1986
<b>Division</b>	A factor with levels E and W indicating player's division at the end of 1986

<b>PutOuts</b>	Number of put outs in 1986
<b>Assists</b>	Number of assists in 1986
<b>Errors</b>	Number of errors in 1986
<b>Salary</b>	1987 annual salary on opening day in thousands of dollars
<b>NewLeague</b>	A factor with levels A and N indicating player's league at the beginning of 1987

From the Fig below we Summarize the dataset in the following conclusion:

1. On an average 380 players had come on bat, and the maximum of them who had come on bat are 600 players. Also, as many as 7 years the players played the major leagues.
2. The maxim of 8 rounds the number of players have batted (RBI).
3. On an average the average salary of the players is roughly estimated to 535.6 thousand dollars. The maximum paid salary was of 2460 thousand dollars.
4. On average 8 times the players have made errors, but its surprising there are as many as 32 errors made.
5. 106.9 times the Assist were provided to these professional players. the maximum number of assist provided is 492 times.

AtBat	Hits	HmRun
Min. : 16.0	Min. : 1	Min. : 0.00
1st Qu.: 255.2	1st Qu.: 64	1st Qu.: 4.00
Median : 379.5	Median : 96	Median : 8.00
Mean : 380.9	Mean : 101	Mean : 10.77
3rd Qu.: 512.0	3rd Qu.: 137	3rd Qu.: 16.00
Max. : 687.0	Max. : 238	Max. : 40.00

Runs	RBI	walks
Min. : 0.00	Min. : 0.00	Min. : 0.00
1st Qu.: 30.25	1st Qu.: 28.00	1st Qu.: 22.00
Median : 48.00	Median : 44.00	Median : 35.00
Mean : 50.91	Mean : 48.03	Mean : 38.74
3rd Qu.: 69.00	3rd Qu.: 64.75	3rd Qu.: 53.00
Max. : 130.00	Max. : 121.00	Max. : 105.00

Years	CAtBat	CHits
Min. : 1.000	Min. : 19.0	Min. : 4.0
1st Qu.: 4.000	1st Qu.: 816.8	1st Qu.: 209.0
Median : 6.000	Median : 1928.0	Median : 508.0
Mean : 7.444	Mean : 2648.7	Mean : 717.6
3rd Qu.: 11.000	3rd Qu.: 3924.2	3rd Qu.: 1059.2
Max. : 24.000	Max. : 14053.0	Max. : 4256.0

CHmRun	CRuns	CRBI
Min. : 0.00	Min. : 1.0	Min. : 0.00
1st Qu.: 14.00	1st Qu.: 100.2	1st Qu.: 88.75
Median : 37.50	Median : 247.0	Median : 220.50
Mean : 69.49	Mean : 358.8	Mean : 330.12
3rd Qu.: 90.00	3rd Qu.: 526.2	3rd Qu.: 426.25
Max. : 548.00	Max. : 2165.0	Max. : 1659.00

Cwalks	League	Division	PutOuts
Min. : 0.00	A:175	E:157	Min. : 0.0
1st Qu.: 67.25	N:147	W:165	1st Qu.: 109.2
Median : 170.50			Median : 212.0
Mean : 260.24			Mean : 288.9
3rd Qu.: 339.25			3rd Qu.: 325.0
Max. : 1566.00			Max. : 1378.0

Assists	Errors	Salary	NewLeague
Min. : 0.0	Min. : 0.00	Min. : 67.5	A:176
1st Qu.: 7.0	1st Qu.: 3.00	1st Qu.: 190.0	N:146
Median : 39.5	Median : 6.00	Median : 425.0	
Mean : 106.9	Mean : 8.04	Mean : 535.9	
3rd Qu.: 166.0	3rd Qu.: 11.00	3rd Qu.: 750.0	
Max. : 492.0	Max. : 32.00	Max. : 2460.0	
		NA's : 59	

The `regsubsets()` method from the 'leaps' package identifies the best model that contains a specified number of predictors, where best is measured using RSS, and conducts best subset selection. The syntax is identical to that of `lm()`. For each model size, the `summary()` command returns the ideal set of variables.

```
Subset selection object
Call: regsubsets.formula(Salary ~ ., data = Hitters, nvmax = 19)
19 Variables (and intercept)
Forced in Forced out
AtBat FALSE FALSE
Hits FALSE FALSE
HmRun FALSE FALSE
Runs FALSE FALSE
RBI FALSE FALSE
Walks FALSE FALSE
Years FALSE FALSE
CatBat FALSE FALSE
CHits FALSE FALSE
CHmRun FALSE FALSE
CRuns FALSE FALSE
CRBI FALSE FALSE
CWalks FALSE FALSE
LeagueN FALSE FALSE
DivisionW FALSE FALSE
PutOuts FALSE FALSE
Assists FALSE FALSE
Errors FALSE FALSE
NewLeagueN FALSE FALSE
1 subsets of each size up to 19
Selection Algorithm: exhaustive
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CatBat
1	( 1 )	" "	" "	" "	" "	" "	" "	" "
2	( 1 )	" "	" "	" "	" "	" "	" "	" "
3	( 1 )	" "	" "	" "	" "	" "	" "	" "
4	( 1 )	" "	" "	" "	" "	" "	" "	" "
5	( 1 )	" "	" "	" "	" "	" "	" "	" "
6	( 1 )	" "	" "	" "	" "	" "	" "	" "
7	( 1 )	" "	" "	" "	" "	" "	" "	" "
8	( 1 )	" "	" "	" "	" "	" "	" "	" "
9	( 1 )	" "	" "	" "	" "	" "	" "	" "
10	( 1 )	" "	" "	" "	" "	" "	" "	" "
11	( 1 )	" "	" "	" "	" "	" "	" "	" "
12	( 1 )	" "	" "	" "	" "	" "	" "	" "
13	( 1 )	" "	" "	" "	" "	" "	" "	" "
14	( 1 )	" "	" "	" "	" "	" "	" "	" "
15	( 1 )	" "	" "	" "	" "	" "	" "	" "
16	( 1 )	" "	" "	" "	" "	" "	" "	" "
17	( 1 )	" "	" "	" "	" "	" "	" "	" "
18	( 1 )	" "	" "	" "	" "	" "	" "	" "
19	( 1 )	" "	" "	" "	" "	" "	" "	" "

	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN	DivisionW
1	( 1 )	" "	" "	" "	" "	" "	" "
2	( 1 )	" "	" "	" "	" "	" "	" "
3	( 1 )	" "	" "	" "	" "	" "	" "
4	( 1 )	" "	" "	" "	" "	" "	" "
5	( 1 )	" "	" "	" "	" "	" "	" "
6	( 1 )	" "	" "	" "	" "	" "	" "
7	( 1 )	" "	" "	" "	" "	" "	" "
8	( 1 )	" "	" "	" "	" "	" "	" "
9	( 1 )	" "	" "	" "	" "	" "	" "
10	( 1 )	" "	" "	" "	" "	" "	" "
11	( 1 )	" "	" "	" "	" "	" "	" "
12	( 1 )	" "	" "	" "	" "	" "	" "
13	( 1 )	" "	" "	" "	" "	" "	" "
14	( 1 )	" "	" "	" "	" "	" "	" "
15	( 1 )	" "	" "	" "	" "	" "	" "
16	( 1 )	" "	" "	" "	" "	" "	" "
17	( 1 )	" "	" "	" "	" "	" "	" "
18	( 1 )	" "	" "	" "	" "	" "	" "
19	( 1 )	" "	" "	" "	" "	" "	" "

	PutOuts	Assists	Errors	NewLeagueN
1	( 1 )	" "	" "	" "
2	( 1 )	" "	" "	" "
3	( 1 )	" "	" "	" "
4	( 1 )	" "	" "	" "
5	( 1 )	" "	" "	" "
6	( 1 )	" "	" "	" "
7	( 1 )	" "	" "	" "
8	( 1 )	" "	" "	" "
9	( 1 )	" "	" "	" "
10	( 1 )	" "	" "	" "
11	( 1 )	" "	" "	" "
12	( 1 )	" "	" "	" "
13	( 1 )	" "	" "	" "
14	( 1 )	" "	" "	" "
15	( 1 )	" "	" "	" "
16	( 1 )	" "	" "	" "
17	( 1 )	" "	" "	" "
18	( 1 )	" "	" "	" "
19	( 1 )	" "	" "	" "

```
[1] "which" "rsq" "rss" "adjr2" "cp" "bic"
[7] "outmat" "obj"
```

**A variable is marked with an asterisk ("\*") if it is present in the associated model.** For instance, this result shows that Hits and CRBI are the only two variables in the optimal two-variable model. Regsubsets() by default only presents results for the top-performing eight-variable model. However, it is possible to return as many variables as needed by using the nvmax option. Here, we fit a model with up to 19 variables.

```
> names(reg.summary)
[1] "which" "rsq"    "rss"    "adjr2"  "cp"    "bic"
[7] "outmat" "obj"
> reg.summary$cp
[1] 104.281319  50.723090  38.693127  27.856220  21.613011
[6]  14.023870  13.128474   7.400719   6.158685   5.009317
[11]   5.874113   7.330766   8.888112  10.481576  12.346193
[16]  14.187546  16.087831  18.011425  20.000000
> reg.summary$adjr2
[1] 0.3188503 0.4208024 0.4450753 0.4672734 0.4808971
[6] 0.4972001 0.5007849 0.5137083 0.5180572 0.5222606
[11] 0.5225706 0.5217245 0.5206736 0.5195431 0.5178661
[16] 0.5162219 0.5144464 0.5126097 0.5106270
> reg.summary$bic
[1] -90.84637 -128.92622 -135.62693 -141.80892 -144.07143
[6] -147.91690 -145.25594 -147.61525 -145.44316 -143.21651
[11] -138.86077 -133.87283 -128.77759 -123.64420 -118.21832
[16] -112.81768 -107.35339 -101.86391  -96.30412
```

By looking at the output below, we can see that the model with 6 variables performs the best overall, according to BIC. There are ten variables in Cp. The adjusted R2 hints that 11 might be the ideal. A model with 5 or less predictors is insufficient, whereas a model with more than 12 predictors is overfitting. Again, no one measure will provide us with an absolutely correct picture.

```
> which.min(reg.summary$cp)
[1] 10
> which.max(reg.summary$adjr2)
[1] 11
> which.min(reg.summary$bic)
[1] 6

> backward = regsubsets(Salary ~ ., data = Hitters, method = "backward")
> reg.summary <- summary(backward)
> reg.summary
```



```

Subset selection object
Call: regsubsets.formula(Salary ~ ., data = Hitters, method = "backward")
19 Variables (and intercept)

```

	Forced in	Forced out
AtBat	FALSE	FALSE
Hits	FALSE	FALSE
HmRun	FALSE	FALSE
Runs	FALSE	FALSE
RBI	FALSE	FALSE
walks	FALSE	FALSE
Years	FALSE	FALSE
CAtBat	FALSE	FALSE
CHits	FALSE	FALSE
CHmRun	FALSE	FALSE
CRuns	FALSE	FALSE
CRBI	FALSE	FALSE
Cwalks	FALSE	FALSE
LeagueN	FALSE	FALSE
DivisionW	FALSE	FALSE
PutOuts	FALSE	FALSE
Assists	FALSE	FALSE
Errors	FALSE	FALSE
NewLeagueN	FALSE	FALSE

		AtBat	Hits	HmRun	Runs	RBI	walks	Years	CAtBat	CHits	CHmRun	CRuns	CRBI
1	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "
2	( 1 )	" "	"*"	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "
3	( 1 )	" "	"*"	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "
4	( 1 )	"*"	"*"	" "	" "	" "	" "	" "	" "	" "	" "	"*"	" "
5	( 1 )	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	"*"	" "
6	( 1 )	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	"*"	" "
7	( 1 )	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	"*"	" "
8	( 1 )	"*"	"*"	" "	" "	" "	"*"	" "	" "	" "	" "	"*"	"*"

		Cwalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
1	( 1 )	" "	" "	" "	" "	" "	" "	" "
2	( 1 )	" "	" "	" "	" "	" "	" "	" "
3	( 1 )	" "	" "	" "	"*"	" "	" "	" "
4	( 1 )	" "	" "	" "	"*"	" "	" "	" "
5	( 1 )	" "	" "	" "	"*"	" "	" "	" "
6	( 1 )	" "	" "	"*"	"*"	" "	" "	" "
7	( 1 )	"*"	" "	"*"	"*"	" "	" "	" "
8	( 1 )	"*"	" "	"*"	"*"	" "	" "	" "

```

> names(reg.summary)
[1] "which" "rsq" "rss" "adjr2" "cp" "bic" "outmat" "obj"
> which.max(reg.summary$adjr2)
[1] 8

```

**The 8 variable model is preferred**, as seen in the result below, according to the Adjusted R2.

## References

1. Bevans ([2022, November 11](#));Datanovia ([2019, December 26](#));
2. *Linear Regression Example in r Using Lm() Function* ([n.d.](#));Zach ([2021, September 29](#));John ([2023, January 25](#));Rithika ([2022, December 29](#))
3. Bevans, R. 2022, November 11. *Hypothesis Testing | a Step-by-Step Guide with Easy Examples*. <https://www.scribbr.com/statistics/hypothesis-testing/>.
4. Datanovia. 2019, December 26. *How to Do a t-Test in r: Calculation and Reporting*. <https://www.datanovia.com/en/lessons/how-to-do-a-t-test-in-r-calculation-and-reporting/>.
5. *Linear Regression Example in r Using Lm() Function*. n.d. <https://www.learnbymarketing.com/tutorials/linear-regression-in-r/>.
6. Rithika, S. 2022, December 29. *Building a Churn Prediction Model on Retail Data Simplified: The Ultimate Guide 101. Learn | Hevo*. <https://hevodata.com/learn/churn-prediction-model/>.
7. Zach, Z. 2021, September 29. *How to Perform Logistic Regression in r (Step-by-Step)*. <https://www.statology.org/logistic-regression-in-r/>.

## Appendix

```
data('mtcars')
head(mtcars)

?mtcars

set.seed(100)
trainIndex <- sort(sample(x = nrow(mtcars), size = nrow(mtcars) * 0.7))
sample_train <- mtcars[trainIndex,]
sample_test <- mtcars[-trainIndex,]
head(sample_train)
head(sample_test)

summary(sample_train)

hist(sample_train$mpg,breaks = 10,xlab="MPG", col = "skyblue", xlim=c(5,35)
)

input<- sample_train
```

```

input$am <- as.factor(input$am)
levels(input$am) <-c("AT", "MT")

table(input$am)

dim(input)

library(ggplot2)
library(caret)
ggplot(input, aes(x=am, y=mpg)) + geom_boxplot(fill="lightgreen")

pairs(mpg ~ ., data = sample_train, col= "red")

options(scipen = 100)
model_step <- step(lm(mpg ~ ., data = mtcars), direction = 'both')
summary(model_step)

step(lm(mpg ~ 1, data = mtcars), direction = 'forward', scope = ~ disp + hp
+ drat + wt + qsec)
model_forward <- lm(formula = mpg ~ wt + hp, data = mtcars)
summary(model_forward)

par(mfrow=c(2,2))
plot(model_step,pch=23,col="orange",cex=2.5,cex.lab=1.6,lwd=3)

fit1 <- lm(formula = mpg ~ wt, data = mtcars)
fit2 <- lm(formula = mpg ~ wt + hp, data = mtcars)
anova(fit1, fit2)

AIC(fit1, fit2)

BIC(fit1, fit2)

library(leaps)
library(ISLR)
library(dplyr)

```

```

summary(Hitters)
Hitters <- Hitters %>% na.omit()

best_subset = regsubsets(Salary ~ ., data = Hitters, nvmax = 19)
reg.summary <- summary(best_subset)
reg.summary
names(reg.summary)

reg.summary$cp
reg.summary$adjr2
reg.summary$bic

which.min(reg.summary$cp)
which.max(reg.summary$adjr2)
which.min(reg.summary$bic)

backward = regsubsets(Salary ~ ., data = Hitters, method = "backward")
reg.summary <- summary(backward)
reg.summary
names(reg.summary)

which.max(reg.summary$adjr2)

## NA

```